

PROJECT SPECIFICATION

Build a Game Playing Agent

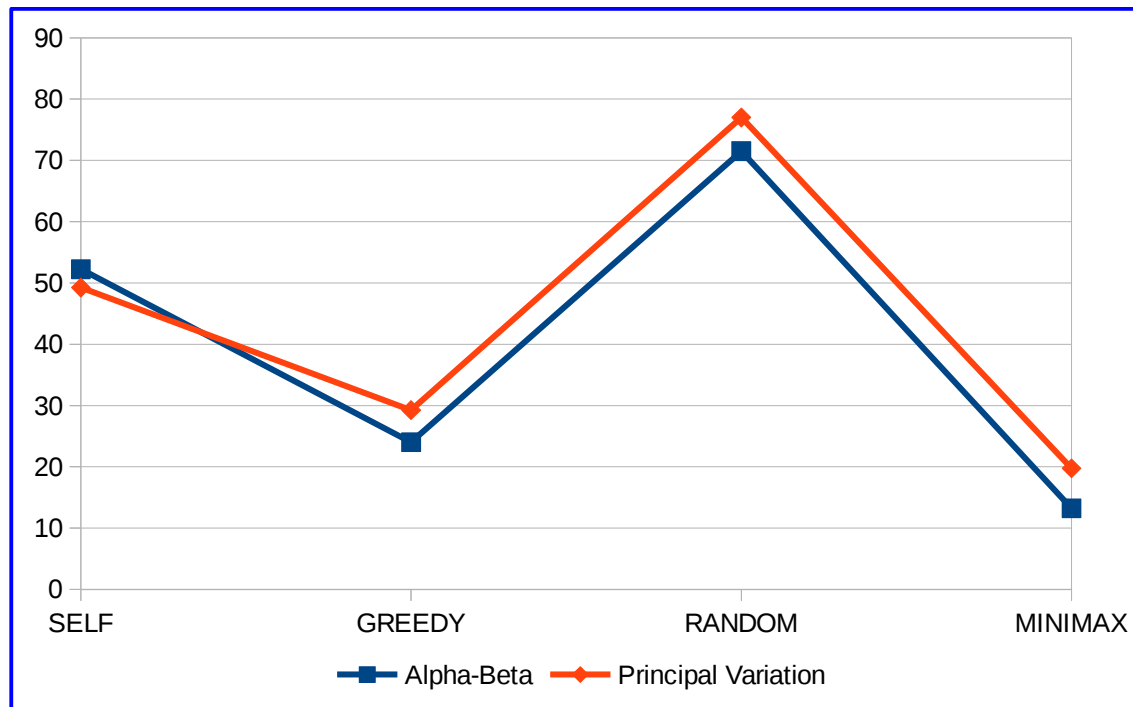
Game Agent Implementation

CRITERIA	MEETS SPECIFICATIONS	STUDENT COMMENTS
<code>.get_action()</code> method calls <code>self.queue.put()</code> at least once before the time limit expires	(AUTOGRADED) Game playing agent can return an action. <ul style="list-style-type: none"><code>.get_action()</code> method calls <code>self.queue.put()</code> at least once before the time limit expires	Implemented in <code>my_custom_player.py</code> (lines 53-54) (Note: Commenting the line depending on whether im running alpha-beta or principal variation)
<code>CustomPlayer</code> successfully plays as both player 1 and player 2 in a full game to a terminal state (i.e., the agent does not deadlock during search, return an invalid action, or raise an exception during a game)	(AUTOGRADED) Game playing agent can play a full game. <ul style="list-style-type: none"><code>CustomPlayer</code> successfully plays as both player 1 and player 2 in a full game to a terminal state (i.e., the agent does not deadlock during search, return an invalid action, or raise an exception during a game)	Implemented in <code>my_custom_player.py</code> All unit tests passing.

Experimental Results

Test Agent	Algotirhm	Run1	Run2	Run3	Run4	Run5	Average	#Games	Fair
SELF	Alpha-Beta	55	47.5	55	57.5	46.2	52.24	40	Yes
SELF	Principal Variation	45	55	46.2	48.8	51.2	49.24	40	Yes
GREEDY	Alpha-Beta	13.8	30	27.5	25	23.8	24.02	40	Yes
GREEDY	Principal Variation	32.5	45	25	28.8	15	29.26	40	Yes
RANDOM	Alpha-Beta	67.5	67.5	77.5	73.8	71.2	71.5	40	Yes
RANDOM	Principal Variation	81.2	77.5	75	72.5	78.8	77	40	Yes
MINIMAX	Alpha-Beta	10	15	17.5	12.5	11.2	13.24	40	Yes
MINIMAX	Principal Variation	15	23.8	20	20	20	19.76	40	Yes

Test Agent	Alpha-Beta	Principal Variation	%Improved
SELF	52.24	49.24	-3.00%
GREEDY	24.02	29.26	5.24%
RANDOM	71.5	77	5.50%
MINIMAX	13.24	19.76	6.52%



Report

CRITERIA	MEETS SPECIFICATIONS	STUDENT COMMENTS
CustomAgent search function uses an advanced search technique	<p><code>CustomAgent</code> class implements at least one of the following:</p> <ul style="list-style-type: none">Custom heuristic (must not be one of the heuristics from lectures, and cannot <i>only</i> be a combination of the number of liberties available to each agent)Opening book (must be at least 4 plies deep)Implements an advanced technique not covered in lecture (e.g., killer heuristic, principle variation search, Monte Carlo tree search, etc.)	<p>I have decided to implement the principle variation search with minimax</p> <p>Implemented in my_custom_player.py</p> <p>I have referred the following pages for implementation along with the lectures and AIMA book</p> <p>https://en.wikipedia.org/wiki/Alpha%E2%80%93beta_pruning</p> <p>https://en.wikipedia.org/wiki/Principal_variation_search</p>
Report includes a table or chart documenting an experiment to evaluate the performance of their agent	<p>Submission includes a table or chart with data from an experiment to evaluate the performance of their agent. The experiment should include an appropriate performance baseline. (Suggested baselines shown below.)</p> <p>Advanced Search Techniques</p>	<p>The results are provided in page 2 of this report and also a chart comparing the baseline (alpha-beta) and the implementation algorithm(principal component search with minimax) is provided. As we can see in the chart, the the agent plays with itself, the results are around 50% and in all the remaining</p>

CRITERIA	MEETS SPECIFICATIONS	STUDENT COMMENTS
	<ul style="list-style-type: none"> Baseline: student must specify an appropriate baseline for comparison (student must decide whether or not <code>fair_matches</code> flag should be used) 	<p>scenarios, the principal variation search works better.</p>
Report answers all required questions	<p>Submission includes a short answer to the applicable questions below. (A short answer should be at least 1-2 sentences at most a small paragraph.)</p> <p>Advanced Search Techniques</p> <ul style="list-style-type: none"> Choose a baseline search algorithm for comparison (for example, alpha-beta search with iterative deepening, etc.). How much performance difference does your agent show compared to the baseline? Why do you think the technique you chose was more (or less) effective than the baseline? 	<ul style="list-style-type: none"> I observe 5.24%, 5.5% and 6.52% improvement with principal variation search for GREEDY, RANDOM and MINIMAX respectively. Please take a look at the table in page 2 of this report. From https://en.wikipedia.org/wiki/Principal_variation_search, we know that principal variation search dominates alpha-beta pruning in the sense that it will never examine a node that

CRITERIA	MEETS SPECIFICATIONS	STUDENT COMMENTS
		can be pruned by alpha-beta; however, it relies on accurate node ordering to capitalize on this advantage